

Distributed (Parallel) Simulation with ns-3

WNS3 2015 Tutorial, Castelldefels (Barcelona), Spain

<https://www.nsnam.org/wiki/AnnualTraining2015>

May 12, 2015

Peter D. Barnes, Jr (LLNL)
Ken Renard (ARL)



Why should you care about distributed (parallel) simulation?


- Faster execution
 - Measure $\sim 10^4$ packet receives/wall clock second/core
- Large models, too big for one compute node
- Heavy-weight nodes
 - DCE applications
 - Virtual machines
 - Core routers with large forwarding tables

Motivation for High Performance, Scalable Network Simulation

- Reduce simulation run-time for large, complex network simulations
 - Complex models require more CPU cycles and memory
 - MANETs, robust radio devices
 - More realistic application-layer models and traffic loading
 - Load balancing among CPUs
 - Potential to enable real-time performance for NS-3 emulation
- Enable larger simulated networks
 - Distribute memory footprint to reduce swap usage
 - Potential to reduce impact of N^2 problems such as global routing
- Allows network researchers to run multiple simulations and collect significant data

ns-3 Execution Scaling

- $10^{~4}$ packets/core/sec
 - Independent of model size
 - 100 cores is 100x faster than 1 core



How many hardware threads do you have?

This laptop

- MacBook Pro, mid 2009
- Intel Core 2 Duo:
2 hardware threads



My other computer

- BlueGene/Q
 - Currently TOP500 #3
 - 8M hardware threads



Outline

- Motivation
- Intro to MPI
- Parallel Discrete Event Simulation
- PDES in ns-3
- Constructing models
- Example
- Error Conditions
- Performance
- Future Capabilities

Big topics. Focus on the concepts and terms needed to understand //ns-3

Motivation

MPI

PDES

ns-3

Models

Example

Errors

Performance

Future


Parallelizing ns-3 Models Is Straightforward, But...

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
HOW MUCH TIME YOU SHAVE OFF	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS

<https://xkcd.com/1205>

And the Reality...




MPI Topics

- Core features
- API specification
- Ranks and communicators
- Point-to-point messages
- Collectives
- Getting and using MPI
- Examples
 - Hello World
 - Simple messaging
 - Ghost cell pattern

Message Passing Interface (MPI)

Features 1

- *De facto* standard programming model for parallel scientific codes (but see Charm++ for an alternative)
- Basic functionality is sending messages (data) between processes, which are called *ranks*
- Core features
 - API specification
 - ~Language independent (FORTRAN, C, C++, Python, Java,...)
 - Supports (doesn't preclude) high performance and scalability
 - Point-to-point (src,dst) messaging, as well as collectives
 - Broadcast, reduction (compute min value), ...
 - Works equally well on distributed and shared memory



<https://computing.llnl.gov/tutorials/mpi/>


MPI Features 2

- API specification, implementation up to “vendor”
 - Vary in performance, runtime launch, ...
 - Architecture-specific libraries can target specialized hardware
 - High-speed interconnects: Infiniband, PAMI, ...
 - Specialized network topologies: Fat-tree, Dragonfly, 5-d torus
 - Specialized network interfaces: low-latency, high-throughput
 - Multi-path routing
 - Multiple implementations can coexist (but not interoperate)
 - OpenMPI, MPICH, IBM, ...
 - Language-specific: mpi4py, mpiJava, ...

MPI Concepts

Ranks and Communicators

- Processes are called *ranks*
- Communicator
 - Group of ranks, numbered $[0, R)$ within the group
 - Enables separating messages by purpose
 - Initial default communicator is `MPI_COMM_WORLD`
 - Several functions for creating communicators, to support specific topologies



MPI Concepts

Point-to-Point Messages

Send a message to a specific rank in a communicator


- `MPI_Send(data, data_length, data_type, destination, tag, communicator)`
 - Data/data_length Message contents
 - data_type MPI-defined data types
Or a custom data type (*i.e.*, a struct)
 - destination Rank Number
 - tag Application tag to distinguish types
 - communicator
- Matching `MPI_Recv()`
- Blocking and non-blocking versions
- Various optimized calls, for managing memory
- Wait for a message, test for new messages

MPI Concepts 3

Collective Communications

Higher level patterns involving more than 2 ranks

- Synchronization
 - MPI_Barrier
- Data movement
 - MPI_Bcast
 - MPI_Scatter
 - MPI_Gather
 - MPI_Allgather
- Reductions
 - MPI_Reduce
 - MPI_Allreduce
- Combinations
 - MPI_Reduce_scatter
 - MPI_Alltoall
 - MPI_Scan



MPI Full API

Environment Management					
MPI_Abort	MPI_Errorhandler_create	MPI_Errorhandler_free	MPI_Errorhandler_get	MPI_Errorhandler_set	MPI_Error_class
MPI_Error_string	MPI_Finalize	MPI_Get_processor_name	MPI_Get_version	MPI_Init	MPI_Initialized
MPI_Wtick	MPI_Wtime				
Point-to-Point Communication					
MPI_Bsend	MPI_Bsend_init	MPI_Buffer_attach	MPI_Buffer_detach	MPI_Cancel	MPI_Get_count
MPI_Get_elements	MPI_Isend	MPI_Iprobe	MPI_Irecv	MPI_Irsend	MPI_Isend
MPI_Issend	MPI_Probe	MPI_Recv	MPI_Recv_init	MPI_Request_free	MPI_Rsend
MPI_Rsend_init	MPI_Ssend	MPI_Ssend_init	MPI_Start	MPI_Startall	MPI_Test
MPI_Test_cancelled	MPI_Testall	MPI_Testany	MPI_Testsome	MPI_Wait	MPI_Waitall
MPI_Waitany	MPI_Waitsome				
Collective Communication					
MPI_Allgather	MPI_Allgatherv	MPI_Allreduce	MPI_Alltoall	MPI_Alltoallv	MPI_Barrier
MPI_Bcast	MPI_Gather	MPI_Gatherv	MPI_Op_create	MPI_Op_free	MPI_Reduce
MPI_Reduce_scatter	MPI_Scan	MPI_Scatter	MPI_Scatterv		
Process Group					
MPI_Group_compare	MPI_Group_difference	MPI_Group_excl	MPI_Group_free	MPI_Group_incl	MPI_Group_intersection
MPI_Group_range_excl	MPI_Group_range_incl	MPI_Group_rank	MPI_Group_size	MPI_Group_translate_ran	MPI_Group_union
Communicators					
MPI_Comm_compare	MPI_Comm_create	MPI_Comm_dup	MPI_Comm_free	MPI_Comm_group	MPI_Comm_rank
MPI_Comm_remote_group	MPI_Comm_remote_size	MPI_Comm_size	MPI_Comm_split	MPI_Comm_test_inter	MPI_Intercomm_create
MPI_Intercomm_merge					
Derived Types					
MPI_Type_commit	MPI_Type_contiguous	MPI_Type_extent	MPI_Type_free	MPI_Type_hindexed	MPI_Type_hvector
MPI_Type_indexed	MPI_Type_lb	MPI_Type_size	MPI_Type_struct	MPI_Type_ub	MPI_Type_vector
Virtual Topology					
MPI_Cart_coords	MPI_Cart_create	MPI_Cart_get	MPI_Cart_map	MPI_Cart_rank	MPI_Cart_shift
MPI_Cart_sub	MPI_Cartdim_get	MPI_Dims_create	MPI_Graph_create	MPI_Graph_get	MPI_Graph_map
MPI_Graph_neighbors	MPI_Graph_neighbors_count	MPI_Graphdims_get	MPI_Topo_test		
Miscellaneous					
MPI_Address	MPI_Attr_delete	MPI_Attr_get	MPI_Attr_put	MPI_Keyval_create	MPI_Keyval_free
MPI_Pack	MPI_Pack_size	MPI_Pcontrol	MPI_Unpack		

Getting and Using MPI

- Check your package manager
- Only the API defined
 - Tool names and configuration vary
 - OpenMPI commands used here for illustration
- Building your code
 - Typically a compiler wrapper script, to ensure correct includes and libraries: `$ mpicc ...`
 - Often hidden inside your build system (`Makefile`, `wscript`)
- Multiple executables
 - Possible to run different executables on different ranks
 - But job launch commands depend on package, so not portable
 - Typically build everything into one executable, select functions based on rank id at runtime

Getting and Using MPI

- Where to run ranks?
 - Single computer: typically defaults to all hardware threads
 - Ad hoc cluster: --hostfile node names and max number of ranks
 - HPC cluster: typically via a batch job system, which selects physical nodes and launches jobs as a shell script
 - “Overcommitment”: running more ranks than cores or hardware threads
- Launching
 - \$ mpirun -n <nrank> <executable>
 - Need to launch on each host

Example OpenMPI Hosts File

```
# This is an example hostfile.  Comments begin with #
#
# The following node is a single processor machine:
foo.example.com

# The following node is a dual-processor machine:
bar.example.com slots=2

# The following node is a quad-processor machine,
# over-subscribing disallowed
yow.example.com slots=4 max-slots=4
```

Example OpenMPI Build and Run

```
$ mpicxx -o hello hello.cc
$ mpirun -np 4 ./hello
Hello World from rank 3 of 4 (35986)
Hello World from rank 0 of 4 (35983)
Hello World from rank 1 of 4 (35984)
Hello World from rank 2 of 4 (35985)
```

Parallel Hello World

- Typical Structure
 1. Include header
 2. Initialize MPI with command-line args
 3. Get world size, my rank index
 4. Parallel code
 - Send messages, synchronize...
 5. Clean shutdown
 6. Build and Launch

hello.cc

```

1 #include <mpi.h>
2
3 int main (int argc, char **argv)
4 {
5   int size, rank, rc;
6   rc = MPI_Init (&argc, &argv);
7   if (rc != MPI_SUCCESS)
8     MPI_Abort(MPI_COMM_WORLD, rc);
9
10  MPI_Comm_size (MPI_COMM_WORLD, &size);
11  MPI_Comm_rank (MPI_COMM_WORLD, &rank);
12
13  printf ("Hello World from rank %d of %d
14        (%d)\n", rank, size, getpid ());
15
16  MPI_Finalize();
17 }
```

Example OpenMPI Build and Run

```

6 $ mpicxx -o hello hello.cc
7 $ mpirun -n 4 ./hello
8 Hello World from rank 3 of 4 (35986)
9 Hello World from rank 0 of 4 (35983)
10 Hello World from rank 1 of 4 (35984)
11 Hello World from rank 2 of 4 (35985)
```

Simple Messaging Example

Rank0 --“Hello”--> Rank1

- Typical Structure
 1. Check #ranks
 2. Message data buffers
 3. Each rank runs different code
 4. Sends paired with Recv
 5. Send and Recv data lengths, types match

Example OpenMPI Build and Run

```
$ mpicxx -o send1 send1.cc
$ mpirun -np 4 ./send1
Rank 1 received message "Hello" (5) from rank 0 tag 0.
```

Parallel Body of send1.cc

```

1 if (size < 2) {
    printf ("Need two ranks\n");
    MPI_Abort(MPI_COMM_WORLD, 0);
}

2 char *msg = (char *)"Hello";
int msg_len = strlen(msg);
char in_msg[msg_len + 1]; // leave space to add null

3 if (rank == 0) {
    int dest = 1;
    rc = MPI_Send (&msg, msg_len, MPI_CHAR, dest,
                  0, MPI_COMM_WORLD);


4 if (rank == 1) {
    int count = 0;
    MPI_Status stat;

    rc = MPI_Recv (&in_msg, msg_len, MPI_CHAR,
                  MPI_ANY_SOURCE, 0, MPI_COMM_WORLD,
                  &stat);
    in_msg[msg_len] = (char) 0;
    MPI_Get_count (&stat, MPI_CHAR, &count);
    printf("Rank %d received message \"%s\" (%d) "
           "from rank %d tag %d.\n",
           rank, in_msg, count,
           stat.MPI_SOURCE, stat.MPI_TAG);
}
}

```

Ghost Cell Design Pattern

- Decomposition
 - Need neighbors' data: stencil
 - Some neighbors are remote
- Solution:
 - Ghosts replicate data
 - Two-phase execution
 - Exchange neighbor data
Communication
 - Compute local update
Computation



Maximize *computation/communication*. Overlap computation and communication.

PDES Topics

- Discrete Event Simulation
 - Mathematical paradigm and time control
 - State and time evolution
 - Event scheduling
 - Time consuming processes
- Parallel DES
 - Logical processors
 - Causality
 - Granted-time synchronization
 - Lookahead
 - Null-message synchronization

Classification of Simulation Techniques


Mathematical Paradigm

Static vs. dynamic time control

Parallelism?

Synchronization Style

Load Distribution



Mathematical Paradigm

Aspect	Discrete	Continuous
Form of model	Discrete systems Automata, agents, particle systems, stochastic processes, etc.	Ordinary or partial differential equations
Time, space, state	Continuous or discrete	All continuous
State changes	Discontinuous in time Constant between state changes	Continuous in time Occasional discontinuities Piecewise differentiable
Mathematical tools	Probability and statistics	Numerical analysis

- Discrete simulation is natural when there are no underlying physical equations


Time Control

Aspect	Event-Driven	Time-Stepped
Event times	Dynamically computed	Statically chosen
Time resolution	(Ideally) Floating point time Zero lower limit on resolution: inherently multi-scale	(Usually) Integer time Nonzero lower limit on resolution
Event distribution in space and time	Sparse and irregular	Dense and regular
Appropriate for	Irregular, asynchronous and/or multi-scale models	Spatially <i>and</i> temporally regular models

- Event-driven execution imposes no timescale
 - Supports simulation with wide dynamic range in natural time scales and/or long quiescent periods


DES State and Time Evolution

- State values change discontinuously in simulation time
 - Constant between state changes
 - Time interval between state changes is not fixed
- Computation (real world time) required to compute new state value
 - Computation occurs at a fixed value of simulation time
- Rate of model evolution not fixed
 - Faster or slower than real time
 - (Best effort) real time, to interoperate with external real systems



DES Event Scheduling

- Objects communicate by sending *messages* = schedule events
 - *Event* is a function call, to be executed Δt in the future—no backwards arrows
 - Typically an event schedules one or more future events
- All event types allowed
 - Event to self
 - Event sends multiple messages
 - Event sends no messages
 - Events can tie
 - Non-FIFO scheduling



Sequential DES Main Event Loop

```
createInitialObjects();
eventList.insert(initialEvents);           // Priority queue on event time

while ( !(terminationCondition() || eventList.empty()) ) do
{
    event e = eventList.removeMinSimTime(); // Choose next event

    simTime = e.getEventTime();            // Set virtual time and unpack event
    object = e.getEventObject();
    method = e.getMethod();
    args = e.getArgs();


    object.method(args);                // Invoke the event method
                                         // May change state of object
                                         // May schedule future events
                                         // May create or destroy objects
                                         // May cancel (delete) future events
}

finalize();
```

Clean separation between *Simulator* and *Application Model*.


Modeling Time-Consuming Process

- Model state values change at an instant in simulation time
 - So how to model time-consuming processes?
- Finite state machine
 - Model state is the FSM state
 - Events cause FSM transitions, schedule future transitions



Parallel Discrete Event Simulation


- Decompose model into *Logical Processes*
 - Separate objects and event queues
 - Execute independently
 - Events for other LPs become messages
 - \sim MPI Ranks



Parallel execution *must* produce exact same results as sequential!


PDES Execution: LPs Advance Independently

- Sometimes ahead in virtual time,
sometimes behind
- More or less real time
per event
- Never backwards!
 - Hallmark of
conservative execution
(Ask me about
optimistic execution ☺)




Need for Synchronization of LPs: Prevent Causality Violations


- Sequential event sequence



- LP 2's perspective:




- Arrival of LP1 event



Need to guarantee no messages arrive in the past (for conservative PDES).


Granted Time Window Synchronization

- If we could guarantee no remote events will arrive before *GrantedTime*
 - All events before *GrantedTime* are safe
 - At *GrantedTime* need to synchronize:
 - Receive and schedule events from other LPs
 - Compute new *GrantedTime*
- Performance
 - Even workload distribution limits cpu idle time
 - Maximize *GrantedTime* to execute more events in parallel between synchronization



Lookahead and LBTS Provide the Granted Time Guarantee


- Model must provide *Lookahead*
 - Minimum delay for remote events
 - Example: network channel link latency + transmission time for smallest packet
- Lower Bound Time Stamp (*LBTS*)
 - Min next event time across all LPs
- $\text{GrantedTime} = \text{LBTS} + \text{Lookahead}$
- Synchronization across LPs is expensive
 - Typically barrier (to wait for slowest LP)
 - Plus (at least one) all gather or reduction
 - Each of these is $\log(N_{LP})$ in time



Finding large *Lookahead* is key to performance

Null-Message Alternative for Static and Sparse LP Graphs

- *GrantedTime* assumes all LPs can message all other LPs
- But my LP graph is sparse. Why synchronize with everyone?
 - Every message sent could communicate my virtual time
 - Guarantee at least one message every *Lookahead*
 - Send *Null-message* when necessary



To Learn More...

- Much of this material from a short course presented spring 2014
 - David Jefferson, LLNL, co-inventor of Optimistic PDES
 - 15 sessions
 - Sequential DES
 - Ties, LBTS, Lookahead
 - Chandy, Misra, Bryant: YAWNS
 - Deadlock
 - Null Messages
 - Dynamic Object Creation
 - Critical Path
 - Speedup
 - Optimistic DES, TimeWarp
 - Global Virtual Time
 - Commitment
 - Checkpointing
 - Rollback and Reverse Computation
 - Dynamic Load Balancing
 - Mixed Discrete and Continuous
 - Slides and videos publicly available:
<http://pdes-course-2014.ucllnl.org>

Parallel ns-3

- History
- PDES in ns-3
- Mechanics
 - Enabling
 - Running
- PDES Simulators
 - GrantedTime
 - NullMessage
- Parallel Models 1
 - The easy way
- Lookahead
- Under the covers:
 - PointToPointRemoteChannel
 - PointToPointNetDevice
- Parallel Models 2
 - The hard way
 - Limitations

Parallel ns-3 History

- Initial release in ns-3.8
 - J. Pelkey and G. Riley, “Distributed Simulation with MPI in ns-3,” WNS3 2011, Barcelona, Spain.
 - Roots from:
 - Parallel/Distributed ns (pdns)
 - Georgia Tech Network Simulator (GTNetS)
- Publications
 - K. Renard, *et al*, “A Performance and Scalability Evaluation of the ns-3 Distributed Scheduler,” SimuTools 2012
 - S. Nikolaev, *et al*, “Performance of Distributed ns-3 Network Simulator,” SimuTools 2013
 - WNS3 2015

PDES in ns-3

Sequential ns-3

- LP is implicit
 - ns3::Simulator
- Event messages
 - Explicit future function calls
 - Schedule (delay, &fn,...)
- Virtual time discipline
 - DefaultSimulatorImpl
 - RealtimeSimulatorImpl
 - VisualSimulatorImpl

Parallel ns-3

- Each rank is an LP
- Event messages
 - Local to LP: explicit future function calls
 - Remote: implicit message send
- Virtual time discipline
 - DistributedSimulatorImpl
 - NullMessageSimulatorImpl
- Lookahead (later)

Enabling Parallel ns-3

- Configure with --enable-mpi
 - Tries to run mpic++
 - Recognizes OpenMPI and MPICH libraries
 - Defines NS3_MPI and either NS3_OPENMPI or NS3_MPICH
- Followed by usual build

Configuring ns-3 With MPI

```
$ ./waf configure --enable-mpi
Setting top to : ...
...
---- Summary of optional NS-3 features:
Build profile      : debug
...
MPI Support      : enabled
...
'configure' finished successfully (1.295s)

$ ./waf build
...
```

Running Parallel ns-3 Scripts

- Waf can't distinguish sequential and parallel
 - Need to specify `mpirun` and number of ranks explicitly

Running Parallel Scripts with waf and mpirun

```
$ ./waf --run simple-distributed
Waf: Entering directory `build/debug'
Waf: Leaving directory `build/debug'
'build' finished successfully (2.118s)
This simulation requires 2 and only 2 logical processors.
Command ['build/debug/src/mpi/examples/ns3-dev-simple-distributed-debug'] exited with code 1

# Multiple ranks on a single computer:
$ ./waf --run simple-distributed --command-template="mpirun -np 2 %s"
Waf: Entering directory `build/debug'
Waf: Leaving directory `build/debug'
'build' finished successfully (2.104s)
At time 1.02264s packet sink received 512 bytes from 10.1.1.1 port 49153 total Rx 512 bytes
At time 1.0235s packet sink received 512 bytes from 10.1.2.1 port 49153 total Rx 512 bytes
At time 1.02437s packet sink received 512 bytes from 10.1.3.1 port 49153 total Rx 512 bytes
At time 1.02524s packet sink received 512 bytes from 10.1.4.1 port 49153 total Rx 512 bytes

# Multiple computers:
$ mpirun -np 2 ./waf -run simple-distributed
```

Switching Between GrantedTime and NullMessage Simulators

- Use environment variable

```
$ NS_GLOBAL_VALUE=\n  "SimulatorImplementationType=ns3::NullMessageSimulatorImpl"\n\n  ./waf --run ...
```

- Use command line:

Selecting the Parallel Simulator from the Command Line

```
bool nullmsg = false;\nCommandLine cmd;\ncmd.AddValue ("nullmsg", "Enable the use of null-message synchronization", nullmsg);\ncmd.Parse (argc,argv);\n\nif(nullmsg) {\n    GlobalValue::Bind ("SimulatorImplementationType",\n                      StringValue ("ns3::NullMessageSimulatorImpl"));\n} else {\n    GlobalValue::Bind ("SimulatorImplementationType",\n                      StringValue ("ns3::DistributedSimulatorImpl"));\n}\nMpInterface::Enable (&argc, &argv);
```


Constructing Distributed Models

The Easy Way

- All ranks construct the full topology
 - All Nodes, NetDevices and Channels
 - Label Nodes with rank: `Node::Node (uint32_t systemId)`
 - All Internet stacks and addresses
 - Good
 - Single code for model construction, runs sequential and parallel
 - Event execution happens in parallel
 - Enables GOD and NIx-vector routing to work
 - Bad
 - Memory is used for nodes/stacks/devices that “belong” to other ranks
(But come to my talk tomorrow ☺)
- Install local applications only
 - Non-local nodes (not on my rank) should not have applications

Where to Get Lookahead?

- Primarily from link latency
- What about shared channels like CSMA or wireless?
 - Latency can be zero
 - Multiple NetDevices
 - *Can't span ranks!*
- Only PointToPoint links can cross ranks
 - Global *Lookahead* is smallest cross-rank latency



Under the Covers: PointToPointHelper::Install

src/point-to-point/helper/point-to-point-helper.cc

```


bool useNormalChannel = true;
Ptr<PointToPointChannel> channel = 0;

if (MpInterface::IsEnabled ()) {
    uint32_t currSystemId = MpInterface::GetSystemId ();
    if (a->GetSystemId () != currSystemId ||
        b->GetSystemId () != currSystemId) {
        useNormalChannel = false;
    }
}
if (useNormalChannel) {
    channel =
        m_channelFactory.Create<PointToPointChannel> ();
} else {
    channel =
        m_remoteChannelFactory.Create<PointToPointChannel> ();


    Ptr<MpReceiver> mpiRecA = CreateObject<MpReceiver> ();
    mpiRecA->SetReceiveCallback (MakeCallback (&PointToPointHelper::OnReceive,
                                              devA->AggregateObject (mpiRecA)));
    // Same for b
}

```

Sequential



Distributed



Under the Covers: Sending a Packet from PointToPointNetDevice

PointToPointNetDevice Call Chain

```
PointToPointNetDevice::Send() {  
    TransmitStart() {  
        PointToPointRemoteChannel::TransmitStart() {  
            MpiInterface::SendPacket();  
        }  
    }  
}
```

- **MpiInterface::SendPacket()**
 - *Packet data*
 - *Receive time* – Local Now() + Latency + Packet Tx duration
 - Remote SystemId (rank)
 - Remote NodeId
 - Remote InterfaceId
- Serialize packet and destination data
- Send to remote rank with non-blocking MPI_Isend()



Under the Covers: Getting a Remote Packet to the PointToPointNetDevice

At end of *GrantedTime*, DistributedSimulatorImpl calls GrantedTimeWindowMPIInterface::ReceiveMessages()

- Reads all pending MPI messages
 - Deserialize target *Receive time*, *NodeId* and *InterfaceId*
 - Deserialize *packet data*
 - Find Node by *NodeId*
 - Find NetDevice on Node with correct *InterfaceId*
 - Get *MpiReceiver* object aggregated to the *NetDevice*
 - *MpiReceiver* merely holds the correct *NetDevice* Callback
 - Schedule *MpiReceiver::Receive* event at *Receive time*

Building a Distributed ns-3 Simulation


- Choose partitioning strategy
 - Label contiguous regions which can't be partitioned
 - CSMA and wireless
 - Select regions which will share a rank
 - Find large point-to-point latencies for good *Lookahead*
 - Minimize communication between ranks
- Build topology as normal, assigning Nodes to ranks
`CreateObject<Node> (rankId)`
- Rewrite topology to improve partitioning
 - CSMA with only 2 nodes
 - Move latency



Constructing Distributed Models

The Hard Way

- Use the ghost cell design pattern to save memory
 - Only create local Nodes, Applications, Internet stacks, NetDevices and Channels
 - Plus “ghost” nodes: remote endpoint of PointToPointRemoteChannel
- Requires *manual intervention*
 - Global and NIX routing do not see entire topology
 - Add static, default routes manually. Hint: IPv6 allows for more “aggregatable” routes
 - Ghost nodes will likely have incorrect remote Nodeld, Interfaceld
 - Must align interface identifiers by hand in same fashion



Limitations of Distributed NS3

- Partitioning is a manual process
- Partitioning is restricted to Point-To-Point links only
 - Partitioning within a wireless network is not supported
 - *Lookahead* is very small and dynamic
- Need full topology in all LPs
 - Exception with careful node ordering, interface numbering, and manual routing

Example

examples/tutorial/third.cc

(These have diverged slightly in ns-3-dev. Differences minimized here.)

1. Include mpi-module.h
2. Same topology, split across Point-to-point link

```

1  *-*- Mode:C++; c-file-style:"gnu"; indent-tabs-mode:nil; -*- */
2
3
4
5
6
7  * This program is distributed in the hope that it will be useful,
8  * but WITHOUT ANY WARRANTY; without even the implied warranty of
9  * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
10 * GNU General Public License for more details.
11 *
12 * You should have received a copy of the GNU General Public License
13 * along with this program; if not, write to the Free Software
14 * Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA
15 */
16
17 #include "ns3/core-module.h"
18 #include "ns3/point-to-point-module.h"
19 #include "ns3/network-module.h"
20 #include "ns3/applications-module.h"
21 #include "ns3/wifi-module.h"
22 #include "ns3/mobility-module.h"
23 #include "ns3/csma-module.h"
24 #include "ns3/internet-module.h"
25
26 // Default Network Topology
27 /
28 / Wifi 10.1.3.0
29 / AP
30 / * * * *
31 / | | | | 10.1.1.0
32 / n5 n6 n7 n0 ----- n1 n2 n3 n4
33 /           point-to-point | | | |
34 /           -----
35 /           LAN 10.1.2.0
36
37 sing namespace ns3;
38
39 S_LOG_COMPONENT_DEFINE ("ThirdScriptExample");
40
41 nt
42 main (int argc, char *argv[])
43 {
44     bool verbose = true;
45     uint32_t nCsma = 3;
46     uint32_t nWifi = 3;
47     bool tracing = false;

```

```

1  *-*- Mode:C++; c-file-style:"gnu"; indent-tabs-mode:nil; -*- */
2
3
4
5
6
7  * This program is distributed in the hope that it will be useful,
8  * but WITHOUT ANY WARRANTY; without even the implied warranty of
9  * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
10 * GNU General Public License for more details.
11 *
12 * You should have received a copy of the GNU General Public License
13 * along with this program; if not, write to the Free Software
14 * Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA
15 */
16
17 #include "ns3/core-module.h"
18 #include "ns3/point-to-point-module.h"
19 #include "ns3/network-module.h"
20 #include "ns3/applications-module.h"
21 #include "ns3/wifi-module.h"
22 #include "ns3/mobility-module.h"
23 #include "ns3/csma-module.h"
24 #include "ns3/internet-module.h"
25
26 #include "ns3/mpi-module.h"
27
28 // Default Network Topology
29 // (same as third.cc from tutorial)
30 // Distributed simulation, split along the p2p link
31 // Number of wifi or csma nodes can be increased up to 250
32 /
33 /           Rank 0 | Rank 1
34 /
35 / Wifi 10.1.3.0
36 / AP
37 / * * * *
38 / | | | | 10.1.1.0
39 / n5 n6 n7 n0 ----- n1 n2 n3 n4
40 /           point-to-point | | | |
41 /           -----
42 /           LAN 10.1.2.0
43
44 sing namespace ns3;
45
46 S_LOG_COMPONENT_DEFINE ("ThirdExampleDistributed");
47

```



Example

examples/tutorial/third.cc

src/mpi/examples/third-distributed.cc

```
25
26
27 1. Different log component name
28
29 2. Command line argument to select N
30
31 // | | | | 10.1.1.0
32 // n5 n6 n7 n0 ----- n1 n2 n3 n4
33 // point-to-point | | | |
34 // =====
35 // LAN 10.1.2.0
36
37 using namespace ns3;
38
39 NS_LOG_COMPONENT_DEFINE ("ThirdScriptExample");
40
41 int
42 main (int argc, char *argv[])
43 {
44     bool verbose = true;
45     uint32_t nCsma = 3;
46     uint32_t nWifi = 3;
47     bool tracing = false;
48
49     CommandLine cmd;
50     cmd.AddValue ("nCsma", "Number of \"extra\" CSMA nodes/devices", nCsma);
51     cmd.AddValue ("nWifi", "Number of wifi STA devices", nWifi);
52     cmd.AddValue ("verbose", "Tell echo applications to log if true", verbose);
53     cmd.AddValue ("tracing", "Enable pcap tracing", tracing);
54
55     cmd.Parse (argc,argv);
56
57     // Check for valid number of csma or wifi nodes
58     // 250 should be enough, otherwise IP addresses
59     // soon become an issue
60     if (nWifi > 250 || nCsma > 250)
61     {
62         std::cout << "Too many wifi or csma nodes, no more than 250 each." << std::endl;
63         return 1;
64     }
65
66     if (verbose)
67     {
68         LogComponentEnable ("UdpEchoClientApplication", LOG_LEVEL_INFO);
69         LogComponentEnable ("UdpEchoServerApplication", LOG_LEVEL_INFO);
70     }
71 }
```

```
// 33 // Rank 0 | Rank 1
34 // 3.0 AP
35 * * 10.1.1.0
36 // n5 n6 n7 n0 ----- n1 n2 n3 n4
37 // point-to-point | | | |
38 // | | | |
39 // LAN 10.1.2.0
40 //
41 //
42 //
43
44 using namespace ns3;
45
46 NS_LOG_COMPONENT_DEFINE ("ThirdExampleDistributed");
47
48 int
49 main (int argc, char *argv[])
50 {
51     bool verbose = true;
52     uint32_t nCsma = 3;
53     uint32_t nWifi = 3;
54     bool tracing = false;
55     bool nullmsg = false;
56
57     CommandLine cmd;
58     cmd.AddValue ("nCsma", "Number of \"extra\" CSMA nodes/devices", nCsma);
59     cmd.AddValue ("nWifi", "Number of wifi STA devices", nWifi);
60     cmd.AddValue ("verbose", "Tell echo applications to log if true", verbose);
61     cmd.AddValue ("tracing", "Enable pcap tracing", tracing);
62     cmd.AddValue ("nullmsg", "Enable the use of null-message synchronization", nullmsg);
63
64     cmd.Parse (argc, argv);
65
66     // Check for valid number of csma or wifi nodes
67     // 250 should be enough, otherwise IP addresses
68     // soon become an issue
69     if (nWifi > 250 || nCsma > 250)
70     {
71         std::cout << "Too many wifi or csma nodes, no more than 250 each." << std::endl;
72         return 1;
73     }
74
75     if (verbose)
76     {
77         LogComponentEnable ("UdpEchoClientApplication", LOG_LEVEL_INFO);
78         LogComponentEnable ("UdpEchoServerApplication", LOG_LEVEL_INFO);
79     }
```



Example

examples/tutorial/third.cc

1. Condition on NS3_MPI
2. Null message selector
3. Initialize MPI
4. Get rank #, number of ranks
5. Check number of ranks
6. Use symbolic names for each rank
7. Create point-to-point nodes

[src/mpi/examples/third-distributed.cc](#)

```
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66 if (verbose)
67 {
68     LogComponentEnable ("UdpEchoClientApplication", LOG_LEVEL_INFO);
69     LogComponentEnable ("UdpEchoServerApplication", LOG_LEVEL_INFO);
70 }
71
72 NodeContainer p2pNodes;
73 p2pNodes.Create (2);
74
75 PointToPointHelper pointToPoint;
76 pointToPoint.SetDeviceAttribute ("DataRate", StringValue ("5Mbps"));
77 pointToPoint.SetChannelAttribute ("Delay", StringValue ("2ms"));
78
79 NetDeviceContainer p2pDevices;
80 p2pDevices = pointToPoint.Install (p2pNodes);
81
82 NodeContainer csmaNodes;
83 csmaNodes.Add (p2pNodes.Get (1));
84 csmaNodes.Create (nCsma);
85
86 CsmaHelper csma;
87 csma.SetChannelAttribute ("DataRate", StringValue ("100Mbps"));
88 csma.SetChannelAttribute ("Delay", TimeValue (NanoSeconds (6560)));
89
90 NetDeviceContainer csmaDevices;
91 csmaDevices = csma.Install (csmaNodes);
92
93 NodeContainer wifiStaNodes;
94 wifiStaNodes.Create (nWifi);
```

```
80 // Sequential fallback values
81 uint32_t systemId = 0;
82 uint32_t systemCount = 1;
83
84 #ifdef NS3_MPI
85
86 // Distributed simulation setup; by default use granted time window algorithm.
87 if(nullmsg)
88 {
89     GlobalValue::Bind ("SimulatorImplementationType",
90                         StringValue ("ns3::NullMessageSimulatorImpl"));
91 }
92 else
93 {
94     GlobalValue::Bind ("SimulatorImplementationType",
95                         StringValue ("ns3::DistributedSimulatorImpl"));
96 }
97
98 MpiInterface::Enable (&argc, &argv);
99
100 systemId = MpiInterface::GetSystemId ();
101 systemCount = MpiInterface::GetSize ();
102
103 // Check for valid distributed parameters.
104 // Must have 2 and only 2 Logical Processors (LPs)
105 if (systemCount != 2)
106 {
107     std::cout << "This simulation requires 2 and only 2 logical processors." <<
108     return 1;
109 }
110
111 #endif // NS3_MPI
112
113 // System id of Wifi side
114 uint32_t systemWifi = 0;
115
116 // System id of CSMA side
117 uint32_t systemCsma = systemCount - 1;
118
119 NodeContainer p2pNodes;
120 Ptr<Node> p2pNode1 = CreateObject<Node> (systemWifi); // Create node with rank 1
121 Ptr<Node> p2pNode2 = CreateObject<Node> (systemCsma); // Create node with rank 2
122 p2pNodes.Add (p2pNode1);
123 p2pNodes.Add (p2pNode2);
124
125 PointToPointHelper pointToPoint;
126
```



Example

examples/tutorial/third.cc

```

66 LogComponentEnable ("UdpEchoClientApplication", LOG_LEVEL_INFO);
67
68
69
70
71
72 PointToPointHelper pointToPoint;
73 pointToPoint.SetDeviceAttribute ("DataRate", StringValue ("5Mbps"));
74 pointToPoint.SetChannelAttribute ("Delay", StringValue ("2ms"));
75
76 NetDeviceContainer p2pDevices;
77 p2pDevices = pointToPoint.Install (p2pNodes);
78
79 NodeContainer csmaNodes;
80 csmaNodes.Add (p2pNodes.Get (1));
81 csmaNodes.Create (nCsma);
82
83 CsmaHelper csma;
84 csma.SetChannelAttribute ("DataRate", StringValue ("100Mbps"));
85 csma.SetChannelAttribute ("Delay", TimeValue (NanoSeconds (6560)));
86
87 NetDeviceContainer csmaDevices;
88 csmaDevices = csma.Install (csmaNodes);
89
90 NodeContainer wifiStaNodes;
91 wifiStaNodes.Create (nWifi);
92
93 NodeContainer wifiApNode = p2pNodes.Get (0);
94
95 YansWifiChannelHelper channel = YansWifiChannelHelper::Default ();
96 YansWifiPhyHelper phy = YansWifiPhyHelper::Default ();
97 phy.SetChannel (channel.Create ());
98
99 WifiHelper wifi = WifiHelper::Default ();
100 wifi.SetRemoteStationManager ("ns3::AarfWifiManager");
101
102 NqosWifiMacHelper mac = NqosWifiMacHelper::Default ();
103
104 Ssid ssid = Ssid ("ns-3-ssid");
105 mac.SetType ("ns3::StaWifiMac",
106 "Ssid", SsidValue (ssid),
107 "ActiveProbing", BooleanValue (false));
108
109 NetDeviceContainer staDevices;
110 staDevices = wifi.Install (phy, mac, wifiStaNodes);
111
112 mac.SetType ("ns3::AnWifiMac");

```

src/mpi/examples/third-distributed.cc

```

119 NodeContainer p2pNodes;
120 Ptr<Node> p2pNode1 = CreateObject<Node> (systemWifi); // Create node with rank 1
121 Ptr<Node> p2pNode2 = CreateObject<Node> (systemCsma); // Create node with rank 1
122 p2pNodes.Add (p2pNode1);
123 p2pNodes.Add (p2pNode2);
124
125 PointToPointHelper pointToPoint;
126 pointToPoint.SetDeviceAttribute ("DataRate", StringValue ("5Mbps"));
127 pointToPoint.SetChannelAttribute ("Delay", StringValue ("2ms"));
128
129 NetDeviceContainer p2pDevices;
130 p2pDevices = pointToPoint.Install (p2pNodes);
131
132 NodeContainer csmaNodes;
133 csmaNodes.Add (p2pNodes.Get (1));
134 csmaNodes.Create (nCsma, systemCsma); // Create csma nodes with rank 1
135
136 CsmaHelper csma;
137 csma.SetChannelAttribute ("DataRate", StringValue ("100Mbps"));
138 csma.SetChannelAttribute ("Delay", TimeValue (NanoSeconds (6560)));
139
140 NetDeviceContainer csmaDevices;
141 csmaDevices = csma.Install (csmaNodes);
142
143 NodeContainer wifiStaNodes;
144 wifiStaNodes.Create (nWifi, systemWifi); // Create wifi nodes with rank 0
145
146 NodeContainer wifiApNode = p2pNodes.Get (0);
147
148 YansWifiChannelHelper channel = YansWifiChannelHelper::Default ();
149 YansWifiPhyHelper phy = YansWifiPhyHelper::Default ();
150 phy.SetChannel (channel.Create ());
151
152 WifiHelper wifi = WifiHelper::Default ();
153 wifi.SetRemoteStationManager ("ns3::AarfWifiManager");
154
155 NqosWifiMacHelper mac = NqosWifiMacHelper::Default ();
156
157 Ssid ssid = Ssid ("ns-3-ssid");
158 mac.SetType ("ns3::StaWifiMac",
159 "Ssid", SsidValue (ssid),
160 "ActiveProbing", BooleanValue (false));
161
162 NetDeviceContainer staDevices;
163 staDevices = wifi.Install (phy, mac, wifiStaNodes);
164
165 mac.SetType ("ns3::AnWifiMac");

```



Example

examples/tutorial/third.cc

src/mpi/examples/third-distributed.cc

```

136 stack.Install (csmaNodes);
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
    address.SetBase ("10.1.1.0", "255.255.255.0");
    Ipv4InterfaceContainer p2pInterfaces;
    p2pInterfaces = address.Assign (p2pDevices);

    address.SetBase ("10.1.2.0", "255.255.255.0");
    Ipv4InterfaceContainer csmaInterfaces;
    csmaInterfaces = address.Assign (csmaDevices);

    address.SetBase ("10.1.3.0", "255.255.255.0");
    address.Assign (staDevices);
    address.Assign (apDevices);

    UdpEchoServerHelper echoServer (9);
    ApplicationContainer serverApps = echoServer.Install (csmaNodes.Get (nCsma));
    serverApps.Start (Seconds (1.0));
    serverApps.Stop (Seconds (10.0));

    UdpEchoClientHelper echoClient (csmaInterfaces.GetAddress (nCsma), 9);
    echoClient.SetAttribute ("MaxPackets", UintegerValue (1));
    echoClient.SetAttribute ("Interval", TimeValue (Seconds (1.0)));
    echoClient.SetAttribute ("PacketSize", UintegerValue (1024));

    ApplicationContainer clientApps =
        echoClient.Install (wifiStaNodes.Get (nWifi - 1));
    clientApps.Start (Seconds (2.0));
    clientApps.Stop (Seconds (10.0));

    Ipv4GlobalRoutingHelper::PopulateRoutingTables ();
    Simulator::Stop (Seconds (10.0));

    if (tracing == true)
    {
        pointToPoint.EnablePcapAll ("third-distributed-wifi");
        phy.EnablePcap ("third-distributed-wifi", apDevices.Get (0));
        csma.EnablePcap ("third-distributed-wifi", csmaDevices.Get (0), true);

        pointToPoint.EnablePcapAll ("third-distributed-csma");
        phy.EnablePcap ("third-distributed-csma", apDevices.Get (0));
        csma.EnablePcap ("third-distributed-csma", csmaDevices.Get (0), true);
    }

```

1

2

2

```

195 address.SetBase ("10.1.1.0", "255.255.255.0");
196 Ipv4InterfaceContainer p2pInterfaces;
197 p2pInterfaces = address.Assign (p2pDevices);

198 address.SetBase ("10.1.2.0", "255.255.255.0");
199 Ipv4InterfaceContainer csmaInterfaces;
200 csmaInterfaces = address.Assign (csmaDevices);

201 address.SetBase ("10.1.3.0", "255.255.255.0");
202 address.Assign (staDevices);
203 address.Assign (apDevices);

204 // If this rank is systemCsma,
205 // it should contain the server application,
206 // since it is on one of the csma nodes
207 if (systemId == systemCsma)
208 {
209     UdpEchoServerHelper echoServer (9);
210
211     ApplicationContainer serverApps = echoServer.Install (csmaNodes.Get (nCsma));
212     serverApps.Start (Seconds (1.0));
213     serverApps.Stop (Seconds (10.0));
214 }

215 // If this rank is systemWifi
216 // it should contain the client application,
217 // since it is on one of the wifi nodes
218 if (systemId == systemWifi)
219 {
220     UdpEchoClientHelper echoClient (csmaInterfaces.GetAddress (nCsma), 9);
221     echoClient.SetAttribute ("MaxPackets", UintegerValue (1));
222     echoClient.SetAttribute ("Interval", TimeValue (Seconds (1.0)));
223     echoClient.SetAttribute ("PacketSize", UintegerValue (1024));

224     ApplicationContainer clientApps =
225         echoClient.Install (wifiStaNodes.Get (nWifi - 1));
226     clientApps.Start (Seconds (2.0));
227     clientApps.Stop (Seconds (10.0));
228 }

229 Ipv4GlobalRoutingHelper::PopulateRoutingTables ();
230 Simulator::Stop (Seconds (10.0));

231 if (tracing == true)
232 {
233     pointToPoint.EnablePcapAll ("third-distributed-wifi");
234     phy.EnablePcap ("third-distributed-wifi", apDevices.Get (0));
235     csma.EnablePcap ("third-distributed-wifi", csmaDevices.Get (0), true);

236     pointToPoint.EnablePcapAll ("third-distributed-csma");
237     phy.EnablePcap ("third-distributed-csma", apDevices.Get (0));
238     csma.EnablePcap ("third-distributed-csma", csmaDevices.Get (0), true);
239 }

```

Example

examples/tutorial/third.cc

src/mpi/examples/third-distributed.cc

1. Enable PCAP tracing on local nodes?
2. Close MPI cleanly

```

223  {
224  UdpEchoClientHelper echoClient (csmaInterfaces.GetAddress (nCsma), 9);
225  echoClient.SetAttribute ("MaxPackets", UintegerValue (1));
226  echoClient.SetAttribute ("Interval", TimeValue (Seconds (1.0)));
227  echoClient.SetAttribute ("PacketSize", UintegerValue (1024));
228
229  ApplicationContainer clientApps =
230  echoClient.Install (wifiStaNodes.Get (nWifi - 1));
231  clientApps.Start (Seconds (2.0));
232  clientApps.Stop (Seconds (10.0));
233 }

234 Ipv4GlobalRoutingHelper::PopulateRoutingTables ();

235 Simulator::Stop (Seconds (10.0));

236 if (tracing == true)
237 {
238     // Depending on the system Id (rank), the pcap information
239     // traced will be different. For example, the ethernet pcap
240     // will be empty for rank0, since these nodes are placed on
241     // on rank 1. All ethernet traffic will take place on rank 1.
242     // Similar differences are seen in the p2p and wireless pcaps.
243     if (systemId == systemWifi)
244     {
245         pointToPoint.EnablePcapAll ("third-distributed-wifi");
246         phy.EnablePcap ("third-distributed-wifi", apDevices.Get (0));
247         csma.EnablePcap ("third-distributed-wifi", csmaDevices.Get (0), true);
248     }
249     else // systemCsma
250     {
251         pointToPoint.EnablePcapAll ("third-distributed-csma");
252         phy.EnablePcap ("third-distributed-csma", apDevices.Get (0));
253         csma.EnablePcap ("third-distributed-csma", csmaDevices.Get (0), true);
254     }
255
256 }
257

258 Simulator::Run ();
259 Simulator::Destroy ();

260 #ifdef NS3_MPI
261     // Exit the MPI execution environment
262     MpiInterface::Disable ();
263 #endif
264
265
266
267
268 return 0;
269

```



Script Output–Identical

```
$ ./waf -run third
```

```
Waf: Entering directory `build/debug'  
Waf: Leaving directory `build/debug'  
'build' finished successfully (2.152s)  
At time 2s client sent 1024 bytes to 10.1.2.4 port 9  
At time 2.01796s server received 1024 bytes from 10.1.3.3 port 49153  
At time 2.01796s server sent 1024 bytes to 10.1.3.3 port 49153  
At time 2.03364s client received 1024 bytes from 10.1.2.4 port 9
```

```
$ ./waf --run third-distributed \  
--command-template="mpirun -n 2 %s --tracing"
```



```
Waf: Entering directory `build/debug'  
Waf: Leaving directory `build/debug'  
'build' finished successfully (2.050s)  
At time 2s client sent 1024 bytes to 10.1.2.4 port 9  
At time 2.01796s server received 1024 bytes from 10.1.3.3 port 49153  
At time 2.01796s server sent 1024 bytes to 10.1.3.3 port 49153  
At time 2.03364s client received 1024 bytes from 10.1.2.4 port 9
```

Cryptic Error Conditions

- Can't use distributed simulator without MPI compiled in
 - Not finding or building with MPI libraries
 - Reconfigure NS-3 and rebuild
- assert failed. cond=" pNode && pMpIRec", file=../src/mpi/model/mpi-interface.cc, line=413
 - Mis-aligned node or interface IDs



Performance Optimizations

- Larger Lookahead
- Synchronization cost grows exponentially with LP count
 - More work per LP is better
 - Speed gains up to 10^{2-3} ranks, depending on model
- Appropriate performance metric
 - Events/sec can be misleading with varying event cost
 - Packet transmissions (or receives) per wall-clock time




Parallel Performance with Large Computation Load: 802.11+OLSR

- Linear scaling out to 128 ranks



Parallel Performance with Small Computation Load: CSMA+Static


- Performance drops at modest number of ranks



Work in Progress

- Automatic memory scaling
 - Automatic ghost nodes, globally unique node IDs
 - (See my talk tomorrow ☺)
- Automatic partitioning, ghost alignment
- Distributed Real Time
 - Versus simultaneous real-time emulations:
 - LP-to-LP messaging gives greater *Lookahead* than independent ns-3 instances connected by emulated network devices
- Scalable default routing
 - AS-like routing between LPs
 - Scalable replacement for GOD or Nix-vector routing with ghost nodes

(Mostly) Parallel Partitioning Tools





Lawrence Livermore
National Laboratory